

# Predicting Price for houses

## Machine Learning - Project 1

### Group 10

```
In [1]: import pandas as pd
import numpy as np
import plotly.express as px
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn import linear_model
```

### Importing Housing Dataset file


```
In [2]: Housing = pd.read_csv("HousingDataSet.csv")
```

```
In [3]: Housing.head()
```

```
Out[3]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0

5 rows × 10 columns

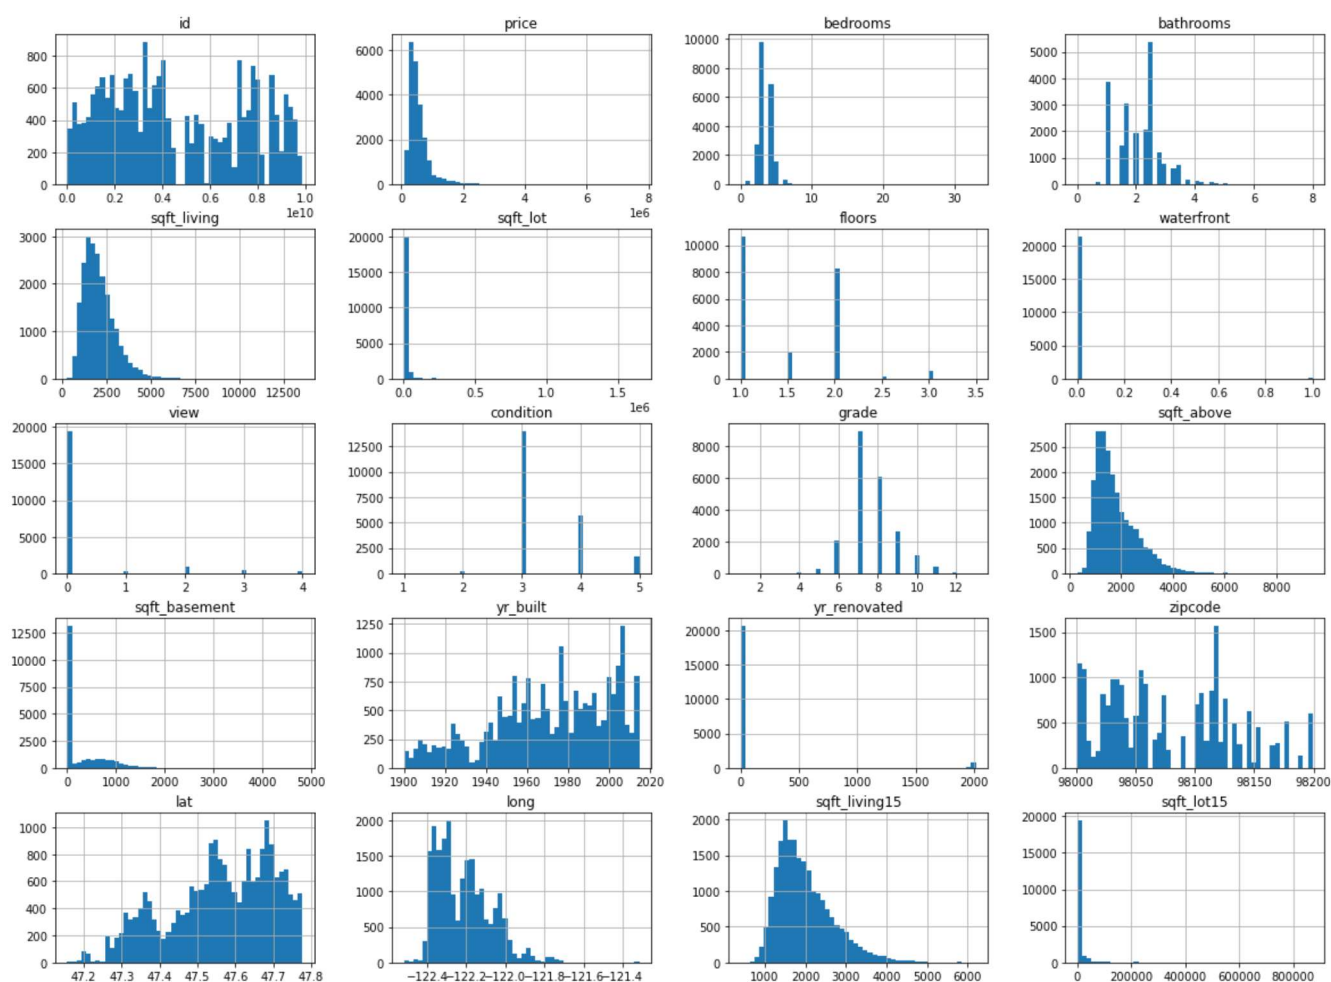


```
In [4]: Housing.shape
Housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21613 non-null  int64
1   date                   21613 non-null  object
2   price                  21613 non-null  float64
3   bedrooms               21613 non-null  int64
4   bathrooms              21613 non-null  float64
5   sqft_living            21613 non-null  int64
6   sqft_lot               21613 non-null  int64
7   floors                 21613 non-null  float64
8   waterfront             21613 non-null  int64
9   view                   21613 non-null  int64
10  condition               21613 non-null  int64
11  grade                   21613 non-null  int64
12  sqft_above              21613 non-null  int64
13  sqft_basement           21613 non-null  int64
14  yr_built                21613 non-null  int64
15  yr_renovated            21613 non-null  int64
16  zipcode                 21613 non-null  int64
17  lat                     21613 non-null  float64
18  long                    21613 non-null  float64
19  sqft_living15           21613 non-null  int64
20  sqft_lot15              21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
In [5]: Housing.hist(bins = 50, figsize =(20,15))
```

```
Out[5]: array([[<AxesSubplot:title={ 'center': 'id' }>,  
  <AxesSubplot:title={ 'center': 'price' }>,  
  <AxesSubplot:title={ 'center': 'bedrooms' }>,  
  <AxesSubplot:title={ 'center': 'bathrooms' }>],  
 [ <AxesSubplot:title={ 'center': 'sqft_living' }>,  
   <AxesSubplot:title={ 'center': 'sqft_lot' }>,  
   <AxesSubplot:title={ 'center': 'floors' }>,  
   <AxesSubplot:title={ 'center': 'waterfront' }>],  
 [ <AxesSubplot:title={ 'center': 'view' }>,  
   <AxesSubplot:title={ 'center': 'condition' }>,  
   <AxesSubplot:title={ 'center': 'grade' }>,  
   <AxesSubplot:title={ 'center': 'sqft_above' }>],  
 [ <AxesSubplot:title={ 'center': 'sqft_basement' }>,  
   <AxesSubplot:title={ 'center': 'yr_built' }>,  
   <AxesSubplot:title={ 'center': 'yr_renovated' }>,  
   <AxesSubplot:title={ 'center': 'zipcode' }>],  
 [ <AxesSubplot:title={ 'center': 'lat' }>,  
   <AxesSubplot:title={ 'center': 'long' }>,  
   <AxesSubplot:title={ 'center': 'sqft_living15' }>,  
   <AxesSubplot:title={ 'center': 'sqft_lot15' }>]], dtype=object)
```



```
In [6]: Housing['bedrooms'].max()
```

```
Out[6]: 33
```


```
In [7]: Corr_Matrix = Housing.corr()
```

```
In [8]: Corr_Matrix['price'].sort_values(ascending = False)
```

```
Out[8]: price            1.000000
sqft_living      0.702035
grade            0.667434
sqft_above       0.605567
sqft_living15    0.585379
bathrooms        0.525138
view             0.397293
sqft_basement    0.323816
bedrooms         0.308350
lat              0.307003
waterfront       0.266369
floors           0.256794
yr_renovated     0.126434
sqft_lot         0.089661
sqft_lot15       0.082447
yr_built         0.054012
condition        0.036362
long             0.021626
id              -0.016762
zipcode          -0.053203
Name: price, dtype: float64
```

```
In [9]: sns.heatmap(Housing.corr())
```

```
Out[9]: <AxesSubplot:>
```




## Selecting the highly correlated features with price for prediction

```
In [11]: X = Housing[['sqft_living', 'bathrooms', 'bedrooms', 'grade', 'floors', 'view']]
y = Housing['price']
```

## Correlation between selected set of variables

```
In [71]: sns.heatmap(X.corr())
```

```
Out[71]: <AxesSubplot:>
```



## Multiple Linear Regression

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
In [14]: lm = LinearRegression()
```

```
In [15]: lm.fit(X_train, y_train)
```

```
Out[15]: LinearRegression()
```

```
In [16]: coeff_df = pd.DataFrame(lm.coef_, X.columns, columns=['Coefficient'])
```

```
In [17]: coeff_df
```

```
Out[17]:
```

	Coefficient
sqft_living	203.161613
bathrooms	-14607.393883
bedrooms	-31084.966289
grade	92334.117869
floors	-21272.496541
view	95120.225965

```
In [18]: predictions = lm.predict(X_test)
```

## Predicting House Price using Multiple Linear Regression

```
In [39]: per_error = 100*(predictions-y_test)/y_test

df_prd_tst = pd.DataFrame({'Predicted Price':predictions.astype('int64'), 'Actual Price':y_test})
df_prd_tst.head(2)
```

```
Out[39]:
```

	Predicted Price	Actual Price	% Error
17384	346470	297000.0	16.656819
722	1399923	1578000.0	-11.284953

## Average Accuracy Score for Multiple Linear Regression

```
In [19]: r2_score = lm.score(X_test,y_test)
print(r2_score*100,'%')
```


```
57.6709013625585 %
```

```
In [20]: sns.distplot((y_test-predictions), bins = 50)
```

C:\Users\sanja\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='price', ylabel='Density'>
```



## KNN Classification Model using K=10

```
In [21]: dataset = Housing[['sqft_living', 'bathrooms', 'bedrooms', 'grade', 'floors', 'view', 'price']]
```

```
In [22]: X = dataset.iloc[:, :6]
y = dataset.iloc[:, 6]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
In [23]: sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.fit_transform(X_test)
```

```
In [24]: classifier = KNeighborsClassifier(n_neighbors=10)
classifier.fit(X_train, y_train)
```

```
Out[24]: KNeighborsClassifier(n_neighbors=10)
```

```
In [25]: y_pred = classifier.predict(X_test)
```

```
In [26]: per_error = 100*(y_pred-y_test)/y_test
Knn10Error = per_error.mean()
```

## Average Accuracy Score for K=10

```
In [27]: print(100+Knn10Error, '%')
```

```
72.03732430437984 %
```

## KNN Classification Model using K=5

```
In [28]: classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(X_train, y_train)
```

```
Out[28]: KNeighborsClassifier()
```

```
In [29]: y_pred = classifier.predict(X_test)
```

```
In [30]: per_error = 100*(y_pred-y_test)/y_test
Knn5Error = per_error.mean()
```

## Average Accuracy Score for K=5

```
In [31]: print(100+Knn5Error, '%')
```

```
75.91960151848336 %
```

## K=5 gives better prediction accuracy than K=10

## Predicting prices of 2 houses using KNN (K=5) and comparing with actual price of house

```
In [32]: per_error = 100*(y_pred-y_test)/y_test
```

```
df_prd_tst = pd.DataFrame({'Predicted Price':y_pred.astype('int64'), 'Actual Price':y_test})
df_prd_tst.head(2)
```

```
Out[32]:
```

	Predicted Price	Actual Price	% Error
--	-----------------	--------------	---------

17384	277500	297000.0	-6.565657
-------	--------	----------	-----------

722	1450000	1578000.0	-8.111534
-----	---------	-----------	-----------

# Lasso Regression Model

```
In [64]: X = Housing[['sqft_living', 'bathrooms', 'bedrooms', 'grade', 'floors', 'view']]
y = Housing['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
clf = linear_model.Lasso(alpha=0.1)
```

```
In [65]: clf.fit(X_train, y_train)
```

```
Out[65]: Lasso(alpha=0.1)
```

```
In [66]: clfpredictions = clf.predict(X_test)
```

## Predicting Price of houses using Lasso Regression Model

```
In [67]: per_error = 100*(clfpredictions-y_test)/y_test

df_prd_tst = pd.DataFrame({'Predicted Price':clfpredictions.astype('int64'), 'Actual Price':y_test})
df_prd_tst.head(2)
```

```
Out[67]:
```

	Predicted Price	Actual Price	% Error
17384	346470	297000.0	16.656819
722	1399923	1578000.0	-11.284953

## Average Accuracy Score for Lasso Regression Model

```
In [68]: predictions = clf.predict(X_test)
r2_score = clf.score(X_test,y_test)
print(r2_score*100,'%')
```

```
57.67089906135141 %
```

**After comparing the accuracy score for all the models used, KNN (K=5) has the highest accuracy, that is 75.92%**

```
In [ ]:
```